

Some technical & scientific challenges I'd like to have working solutions for

Fred “pappy” Raynal
fraynal@quarkslab.com
@fredraynal



Who am I?

- Born: 17th March 1972
- 1st computer around 10: a ZX81 with 1KB of RAM
- Bac C in 1990 (now called S)
- Engineering degree + Military service + DEA (Master) in 1996-1997
- PhD. in 2001 at INRIA (wrote articles for LinuxMag at the same time)
- Started MISC, a French security related magazine, in Jan. 2002 - stopped 13 years later!
- Started SSTIC, a French security conference, in Jun. 2002 - stopped 4 years later!
- Co-created and co-run EADS CCR (IW) until 2006
- Created and run Sogeti ESEC R&D until 2011
- CEO and founder of Quarkslab since then

Preamble

- Some technical challenges I'd like to have **working** solutions for
- Working does not mean:
 - 100% efficient or reliable (f*** corner cases)
 - I don't care for technical / scientific elegance
- Working means:
 - Scalable: yes, size matters
 - Answering at least some parts of the problem

Roadmap

- Now classic (but not working):
 - IR, malware analysis & classification
 - Vulnerability: research, exploitation, ...
- What I also need:
 - Obfuscation
 - Patching

A close-up photograph of a man with dark hair and a beard, wearing a white tank top. He is looking upwards and to his right with an intense expression. He is holding a black microphone in his right hand, which is positioned near his chest. His left arm is extended outwards and slightly upwards, with his hand pointing towards the upper right corner. The background is dark, making the subject stand out.

Friends will be friends:
Incident response

What is Incident Response?

- Incident response is an organized approach to addressing and managing the aftermath of a security breach or attack (also known as an incident).
- The goal is to handle the situation in a way that limits damage and reduces recovery time and costs.

Why do we care?

- Breaches are caused by exploits and “malware”
 - rootkit, spear phishing, virus, APT, ...
- So many different kinds of “malware”, we need to automatize
 - Attack noise: random attacks not often advanced ...
 - but could be used to hide a more serious threat
 - “Targeted” attacks: attacks crafted for a specific target
 - Does not mean attack vectors are exclusively used

Incident response in 2 questions

- Is this a threat?
- If yes, how to fix it?

Dealing with incident response

- A 4-stage process:
 - Init: setup tools and knowledge base (learning)
 - Update: keep learning with new binaries / processes but real-time now
 - Alert: the questions one asks when a threat is detected
 - Remediate: fix and prevent

The false problem: unpacking

- Packing: yes, malware can be packed, but unpacking is easy
 - Why unpack? To be able to perform static analysis
 - Simple tricks are good enough to unpack most of the malware
 - unpacking library
 - kernel instrumentation, non public hypervisor

The problem nobody cares about: *morphism

- Less than 2% of malware use polymorphism
- None known to use metamorphism
- It does not mean it won't one day, but we already have enough sh*t to deal with “basic” and “advanced” threats
 - The day it will be useful, analyzis will be made manually at first

Detection?

- Detection is not a problem anymore => rely on many unreliable tools
 - Specific tools: AV, sandbox, hash database
 - Side channels: big data (DNS, proxy, url, ...), behavioural analysis
- Accept you are going to be compromised by “something” “some day”
 - Be ready to fix quick and clean

Stage #1: init

- Classify **files** (not only binaries)
 - 3 classes: clean, malware, unknown
- Howto:
 - Human analysis: expensive and super-slow
 - **Similarities**: does it really scale?
- The KISS trick:
 - Use a hash database of the users and servers masters

Stage #2: update

- A new file arrives, how to classify it?
 - Clean, malware, unknown
- Usually 2 strategies:
 - Static: when the file is created
 - Dynamic: when the file is executed
- Scaling again: mail, software updates, download, ...

Stage #3: alert

- Quickly needs answer to:
 - Q1: Where is the threat?
 - Can use IOC for that, if one has a way to run that on all the attack surface

Stage #4: remediate

- Quickly needs answer to:
 - Q2: How to clean it?
 - Can I “clean” the threat or should I reinstall everything?
 - From a static / dynamic analyzis, can I provide a way to clean the malware?
 - Learn from attacks: **create new signature / IOC / ***
 - Anything helping to detect this attack again

Conclusion

- Dealing with **many** files: don't work only on malware samples
 - Legitimate, created on the fly, malware, documents, ...
- Dealing with **huge** binaries: use Acrobat Reader as the new Lena
 - Many functions, obviously even more basic blocks
- Any solution not scaling with these 2 constraints can not apply
 - Yes, I prefer to have less matching but to have it working on my file stream



Another one bites the
dust:
vuln & exploits

Vulnerability: from bug to exploit

- Every software has bug(s)
- Bugs might have consequences:
 - DoS: a simple crash, memory exhaustion, ...
 - Code exec: change the legitimate control flow of the target program
- Won't explain why we care for bugs & exploits

Bugs & Exploits

WARNING

Bug research != Exploit development

Strategies to find bugs: hand-made

- Hand-made: sources | binary analyzis
 - Allows to find the best bugs ...
 - Including the logic ones
 - ... but so time consuming
 - how to: top / down, bottom / up, bug types

Strategies to find bugs: half-automatic

- Half-automatic: help to locate points of interest
 - Diffing: if there was a bug there before, another one might still be around
 - Good opportunities vs. difficult context analyzis
 - Source analyzis: look for vuln in a code
 - Soundness vs relevance issues + hard to setup
 - Fuzzing: send malformed inputs to a program
 - KISS vs difficult context analyzis

Strategies to find bugs: automatic

- Automatic: mostly abstract interpretation to prove code
 - AI proves a class of bug is absent in a code
 - But if it fails, is it because:
 - An over-approximation or a real bug?
 - Because the code does not comply to the prover
 - Not relevant for bug finding because:
 - Barely scales on “real” code (embedded / critical systems only)
 - Almost impossible to use without 3 PhDs

Exploitation

- No more automatic methods to exploit
- Don't necessarily need a detailed analysis of the root cause
- Need to understand the memory map:
 - Local variables, registers, internal structures
 - Controlled memory or not
 - Context execution: ASLR, DEP, stack cookie, ...
- Today's exploits are craftsman work

Source audit is the old
black

Source code tools

- Cons
 - Many are VERY expensive
 - Requires to have 3 PhDs to be able to use one
 - Relies on non-trendy language (O-Caml)
- Pros
 - Would love to find some but never used one of these tools ...

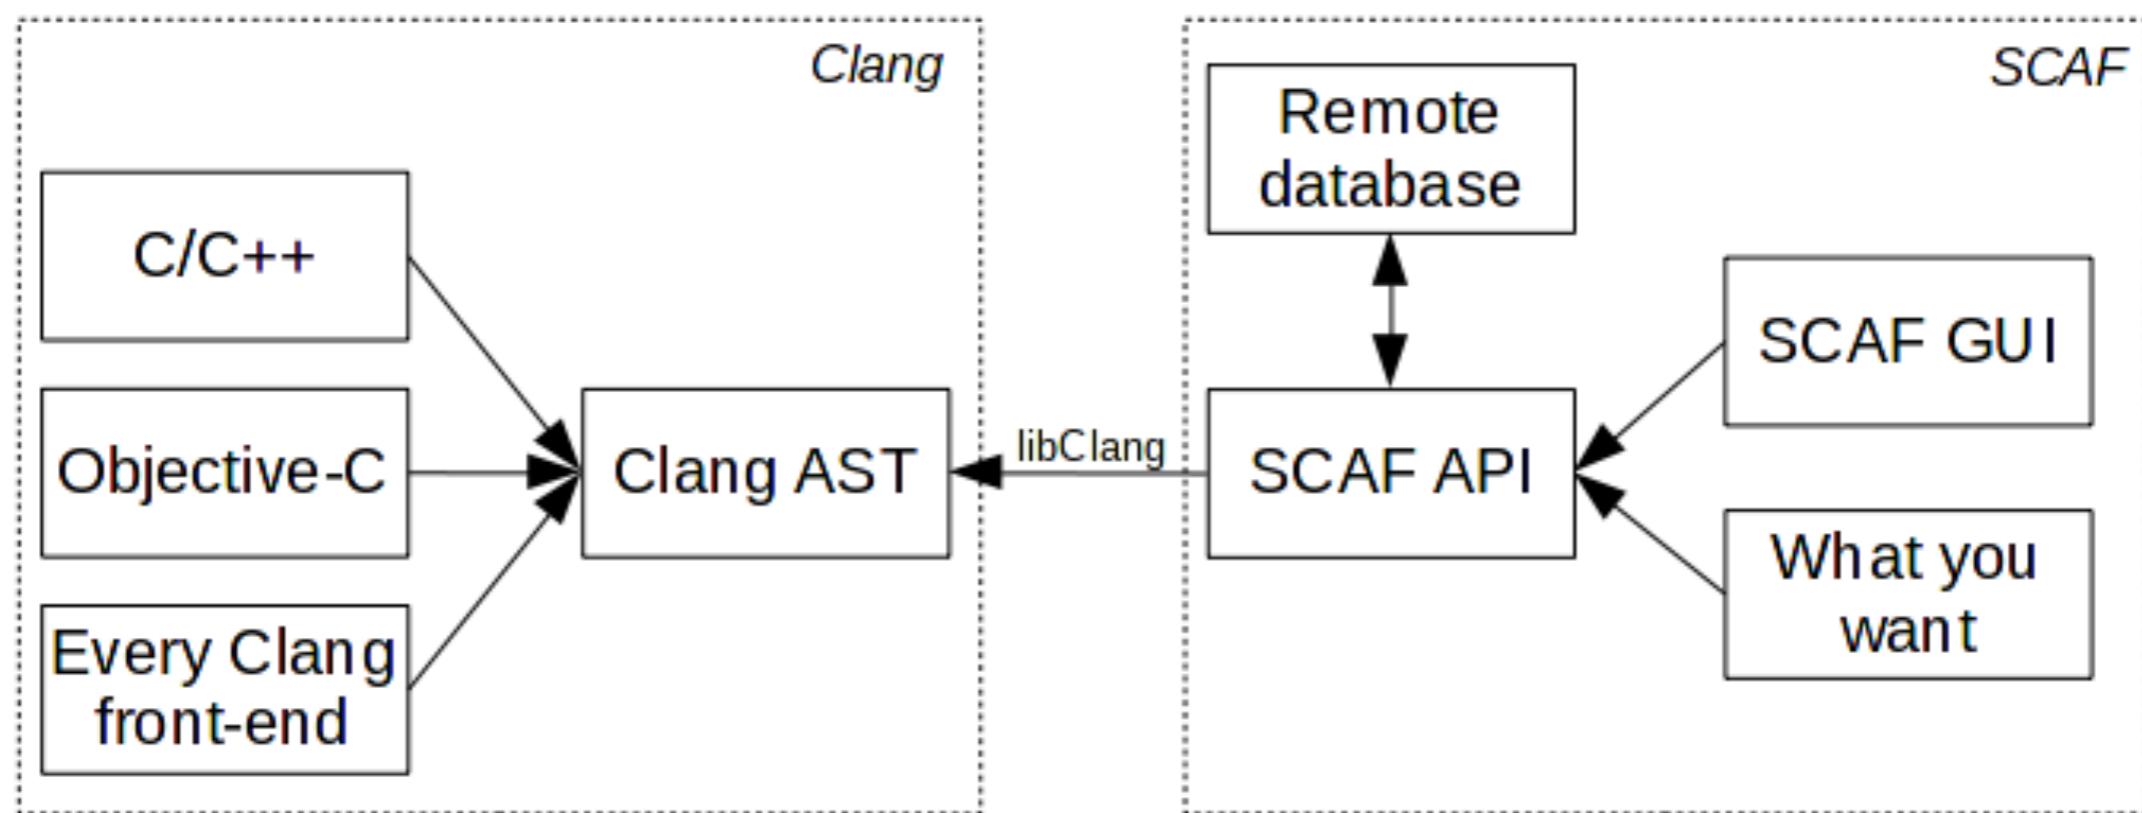
Handmade source audit

- What to look for: kind of signatures
 - Functions (ex.: memory handling), pointer arithmetic, bug class (integer overflow), ...
- How to look for
 - Top / down : start from the main() and move down to the leaves
 - Bottom / up : start from a leaf and get backward to the main()
 - Focused: a specific class (ex.: strings), feature (ex.: a parser)

Our tool^w PoC: SCAF

- Wanted killer features
 - Contextualization: declaration, xref, CFG, ...
 - Framework: want to make plugins and automatize sometimes
 - Collaborative: several users at the same time
 - Scalable: test vector was the linux kernel

SCAF in one picture



Source analysis

- Important difference
 - Usual tools: look for everything
 - Human analyst: look for one thing
- I don't want tools trying to answer questions I did not even ask
- Good test cases: Linux kernel, openssl, Firefox, image libraries...

Clash of clans

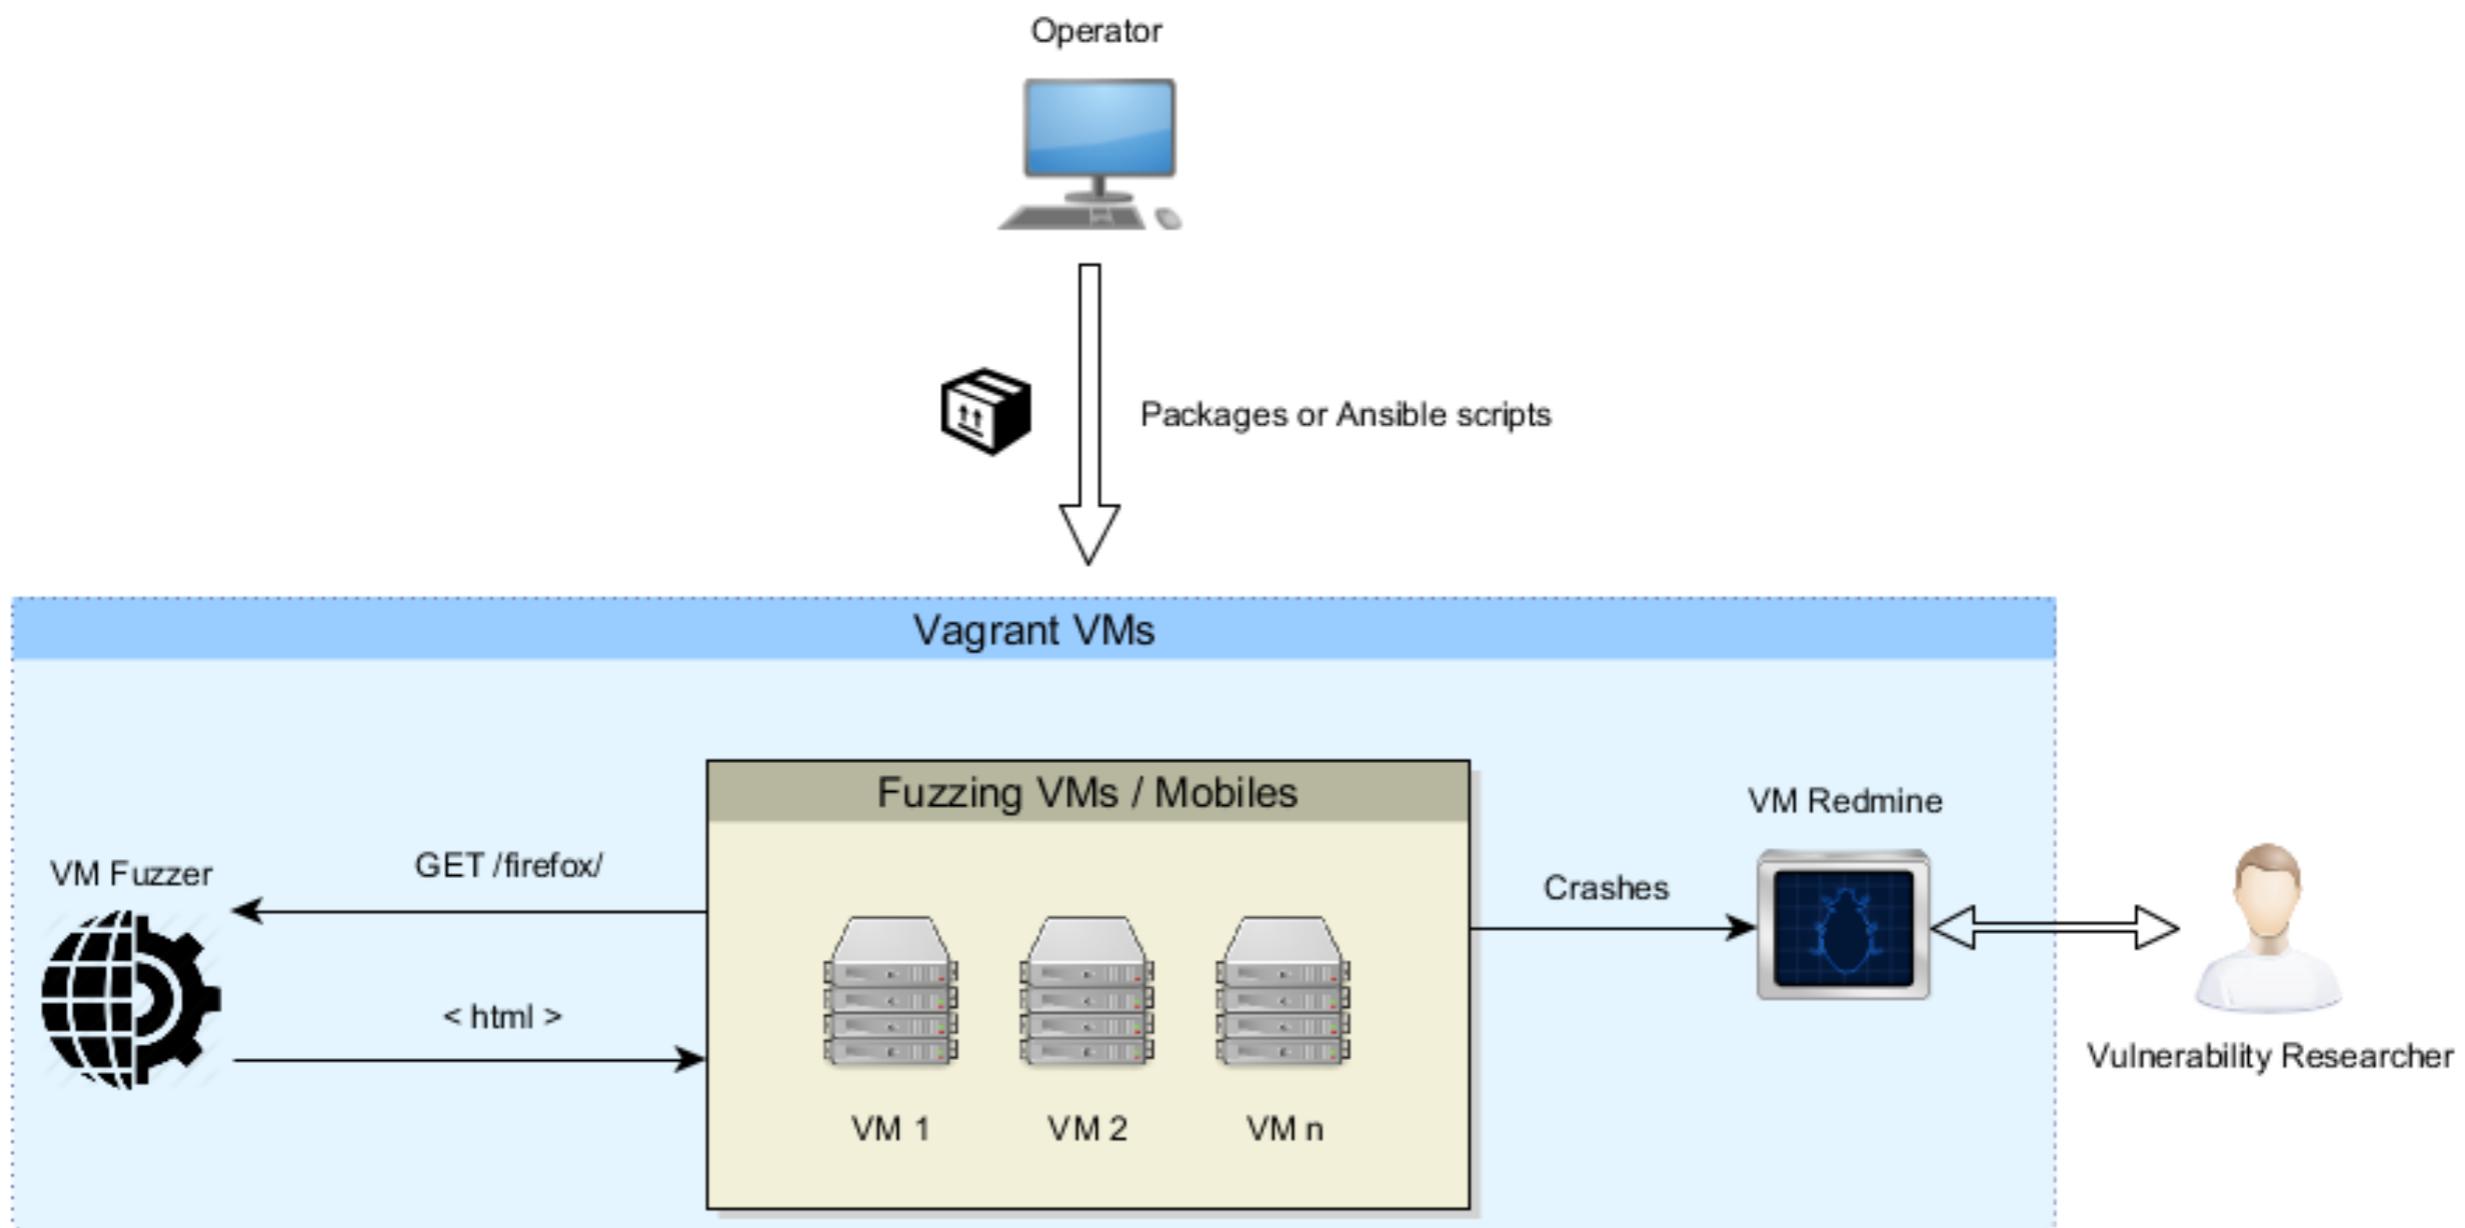
- “Hacker” world: efficiency first, barely strong theory, deal is the current problem the guy is working on
- “Academic”: unusable by a normal human (theory), not user friendly (UI), works on small examples (scalability)
- Conclusion: still a lot of research to do in source analyzis, but should focus on making the current results accessible / usable
- Bonus: would also love to improve context based on execution traces

Fuzzing is the new
black

Fuzzing 101

- Definition (fuzzing): transform sh*it into gold
- Can start very cheap up to a very complex infrastructure
- The best approach today because the best ratio effort / gain

Complex fuzzing infra



(Dumb) Fuzzing 102

- Automatic:
 - Mutation: transform an input into another
 - Agent / supervisor / monitor : keep an eye of the target's behaviour when it is feed with the input
 - Triage: keep record of “interesting” events
- Hand made:
 - Crash reduction: reduce sample to its minimum
 - Crash analyzis: is it exploitable or not?

Problems with fuzzing?

- Code coverage: try to reach every basic block
 - Easy to measure at execution but slows down the execution
 - Greatly depends on non mutated inputs (corpus)
- Data coverage: try to cover the data space
 - Attempt to solve the problem with symbolic / concolic execution to find / solve the constraints
 - Not reliable and super slow
- “Interesting” event: what is a deviant behaviour?

Code coverage at scale

- Corpus distillation
 - Gather a large set of inputs (the corpus)
 - For each, generate an execution trace to store
 - Select the minimal subset in the input covering all the code
 - Start usual dumb fuzzing with this subset as to-be-fuzzed corpus

Corpus distillation numbers

- 20 TB of SWF files
- 1 week of computation on 2000 CPUs to select the minimal corpus => 20 000 files
- 3 weeks of mutation / execution on the 2000 CPUs => 400 unique crashes
- 106 security bugs acknowledged by Adobe
- Source: <http://googleonlinesecurity.blogspot.fr/2011/08/fuzzing-at-scale.html>

Corpus distillation issues

- Mostly only Google is able to do that because they already have all the files
- Does not really solve the problem, use scale for that ... but at least, **it works!**
 - Also means there is a big research to do there ;)

Afl - American Fuzzer Loop

- Based on edges, not basic blocks anymore
- Focus on transitions rather than instructions
- Target is instrumented either at compilation or execution
- Manage its own corpus through genetic algorithms
- Provide tools to manage crashes, corpus, *
- Best fuzzer today: <http://lcamtuf.coredump.cx/afl/>

Challenges with fuzzing

- Corpus: have minimal samples with max code | data coverage
- Code coverage: block vs. edge? Static with use of subgraphs?
- Data coverage: symbolic, concolic
- Crash reduction
- Exploitability

The march of the black queen: Code obfuscation



Obfuscation

- Obfuscation: make programs difficult to analyze
 - Must protect both code and data sometimes
 - Focus very often on the CFG only, not on the data
- Constraint: try not to kill the performances (CPU, memory)

Some techniques

- Opaque predicates, code flattening, inlining / outlining, virtual machines...
- Mixed boolean-arithmetic expressions, string encryption, parameters or prototypes unification...
- Can be source 2 source, source 2 bin, bin 2 bin, ...

Obfuscation & research

- Obfuscation is a very secretive area
 - Pretty much no publications, some very obscure patents
 - You have to sign 20 NDAs when working on attacking obfuscation
 - Players: Google, MS, Adobe, Apple vs. Arxan, Cloakware, Safenet, ...
 - Also some home-made obfuscation (remember the “dont do crypto if you are not a cryptograph” ?)
- I can just tell that we broke the crypto in most of the DRM we attacked

Don't solve the wrong problem

- Problem #1: analysis
 - Exactly the same constraints as for malware or vuln research, no need of a specific focus
- Problem #2: “difficult”
 - Define what is “difficult” to analyze...
 - Assessment always performed on the binary

Challenge: get a complexity measure

- Source or binary analysis keeps saying “program is difficult to analyze”: use that as a metric for obfuscation?
- Metrics on the CFG, the diversity / similarity in a code (to avoid pattern matching)
- What about the data flow?
 - Today: MBA or float are used because no tool supports them properly



Obfuscation REALLY needs to have a good
measure for complexity

Bonus: obfu as crypto



- Apply Kerckhoffs principle to obfuscation's design

We need a strong obfuscation that can “resist” reverse while every detail is public

A man with dark hair and a serious expression, wearing a white, flowing robe. He has a large, irregular tear or patch on his right shoulder, revealing a red, textured material underneath. His left arm is extended forward, palm up, as if presenting something.

Crazy little thing called
patch

Patching as a challenge

- Differing code is a well-known way to find vuln
- Differing binaries has been used a lot to find vuln
- But how to patch once a vuln is known?
- Even better: how to patch an unknown vuln?

When patching makes a difference

- Desktop: most are relying on automatic updates
- Mobile: automatic updates will come there too
 - iOS does it already, Android will anytime soon
- Embedded | critical systems: ...
 - IoT: 2 billions devices to patch at once?
 - Critical systems: some can reboot only once a year
- Abandonware: yes, it does exist also for important systems

Dynamic patching: memory access

- White list: every access is forbidden, it is granted only once it is proven to be safe
 - Very generic, can patch even unknown bugs
- Black list: every access is granted except where a bug was found (e.g. by fuzzing) and a bit around
 - Bug specific but avoid to analyze the root cause

Dynamic patching: input validation

- Determine the list of inputs for a program
 - From network, files, user, ...
- Wrap all the inputs with a function call before initialisation
 - Can use white or black list approaches

One patch to rule them all

- Applying the patch
 - Statically: modifying the binary :/
 - Dynamic: inject in the process, load a library, ...

Additional challenges

- Performance: the patch must have a minimal impact on the target program
- Soundness: the program must work in the same way (with at least a bug removed): can we “prove” it?
 - Re-run with the bug triggering sample
 - Try with many other samples

QUEEN

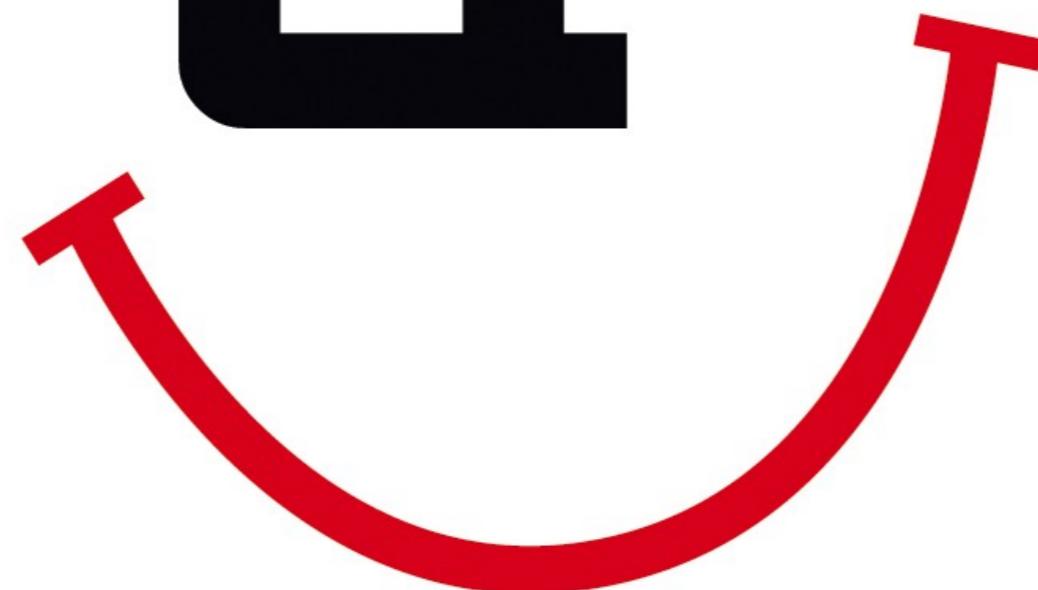


THE SHOW MUST GO ON

Conclusion

- There are still a lot of challenges
- Please, think **big** first, then **sound**
- TOP 3 :
 - Similarities based on many big files
 - Make code analysis usable
 - Patching unknown vuln
- Bonus: think different: firmware, complete system, bootloader...

Questions?



Fred “pappy” Raynal
fraynal@quarkslab.com
@fredraynal