

Attack on Titan M

Vulnerability Research on a Modern Security Chip

Damiano Melotti
Maxime Rossi Bellom

What is Titan M?

- Security chip made by Google,
for Pixel devices
- Implements critical security features
 - Keymaster/Strongbox, Weaver, AVB, etc
- Client-server model
- Introduced for:
 - Mitigate side-channel attacks
 - Protect against hardware tampering

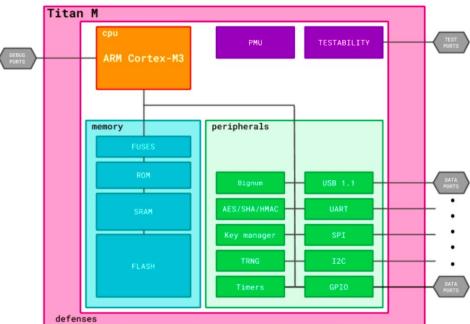


Our previous work in 4 slides

Specification

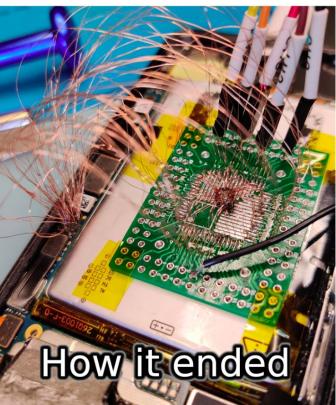
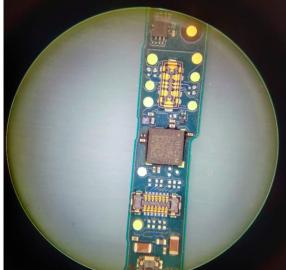
Hardened SoC based on ARM Cortex-M3

- Anti-tampering defenses
- Cryptographic accelerators & True Random Number Generator
- UART for logs and console
- SPI to communicate with Android



Hardware Reverse: Finding SPI

How it started



How it ended



Firmware Security Measures

- Secure boot (images are signed and verified at boot)
- No MMU, but MPU to give permissions to the memory partitions
- Only software protection: hardcoded stack canary checked in the SVC handler

```
if (*CURRENT_TASK->stack != 0xdeadd00d) {  
    next = (int)&CURRENT_TASK[-0x411].MPU_RASR_value >> 6;  
    log("\n\nStack overflow in %s task!\n", (&TASK_NAMES)[next]);  
    software_panic(0xdead6661, next);  
}
```



What can we do with the exploit?

Vulnerable buffer placed just before

- runtime data of the chip...
- ... and the list of command handler pointers

→ overwrite command handler addresses
to gain code execution!



What we already did

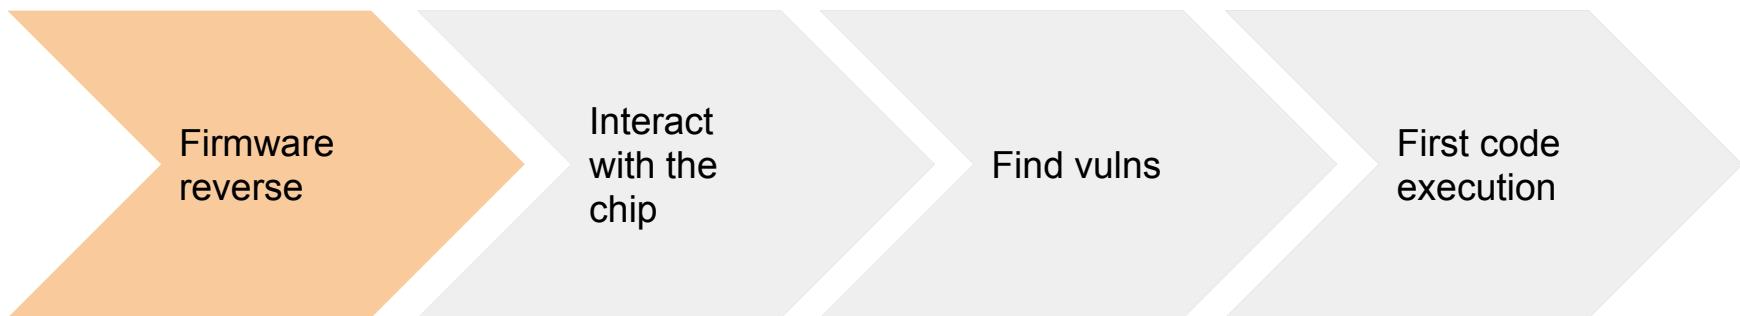
Firmware
reverse

Interact
with the
chip

Find vulns

First code
execution

What we already did

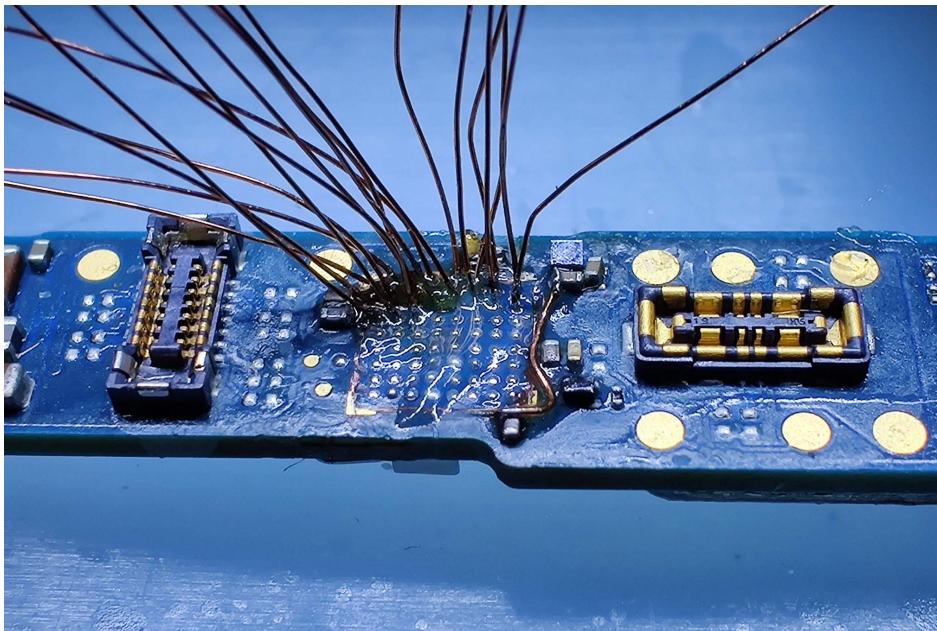


What we already did



What we already did

Implemented some tools to interact with the chip

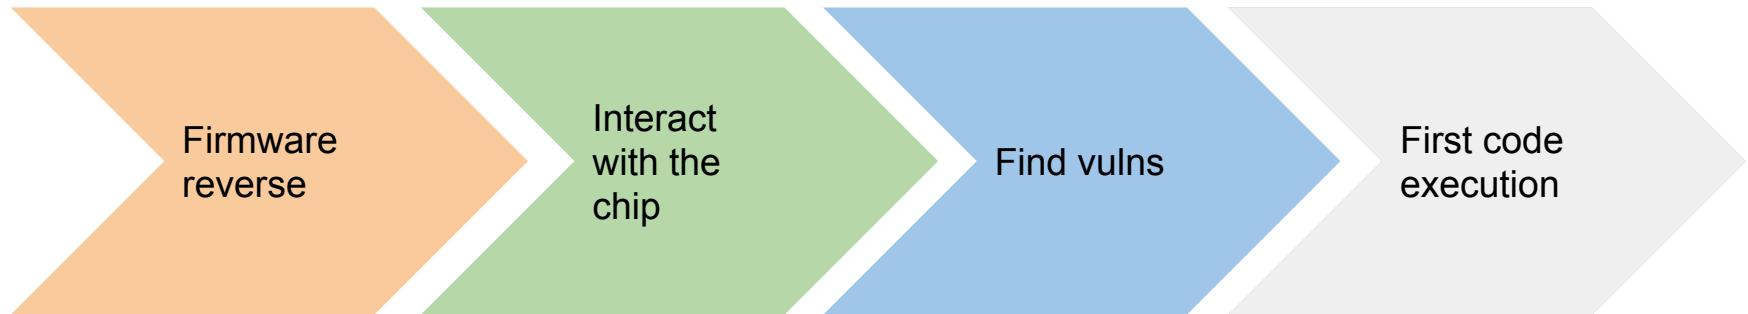


Sniff and send custom commands

- From Android using Frida and our tool **nosclient**
- On this hardware level thanks to @doegox's magic hands

What we already did

Several vulnerabilities reported, including a downgrade issue



Published the first code execution on this target

What we already did

Several vulnerabilities reported, including a downgrade issue



- CVE-2021-0939: A memory leak allowing to reveal parts of the Boot ROM
- CVE-2021-1043: A downgrade issue allowing to flash any firmware
 - With a side effect: all the secrets are erased

What we already did

Leak various hidden parts of the firmware



Including the Boot ROM

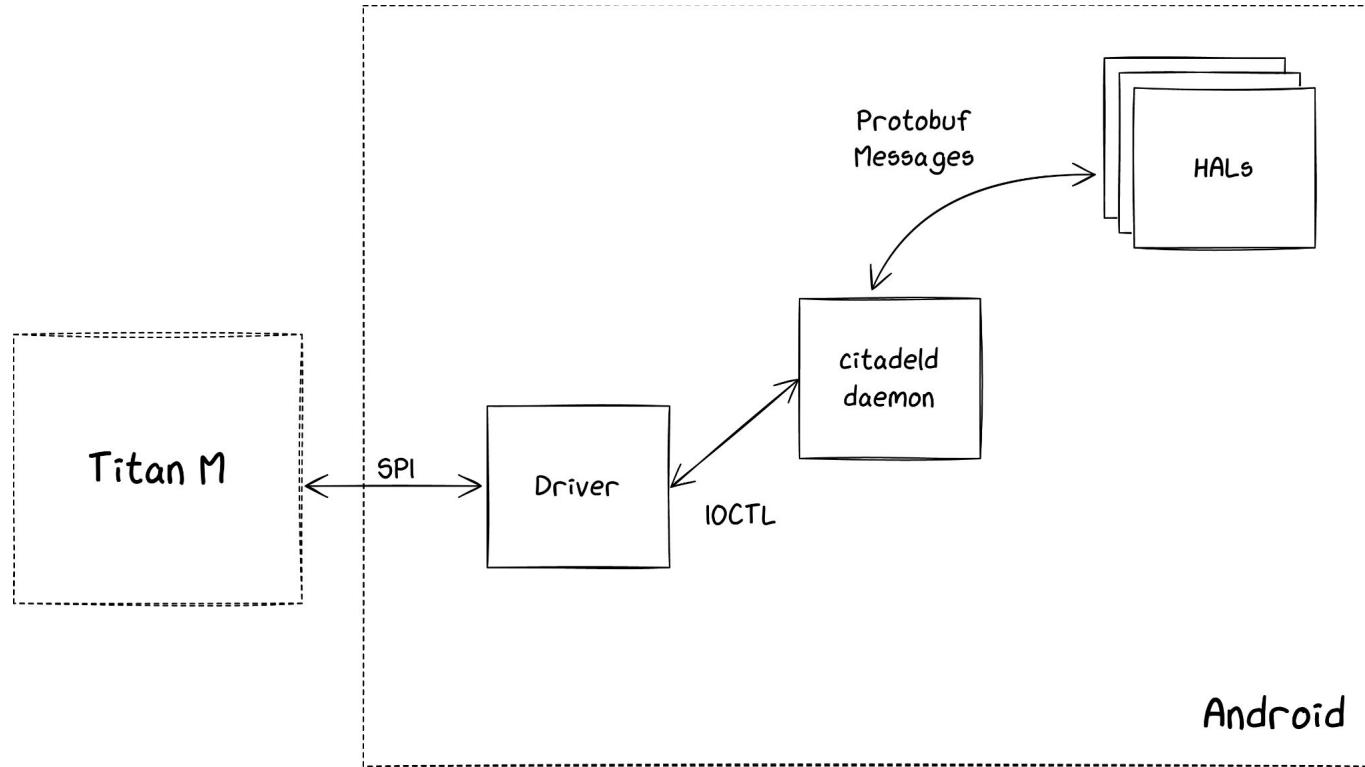
TL;DR: what we learned

- Security chip based on ARM Cortex-M3
- Closed source but based on EC
 - An open source OS made by Google
 - Written in C and conceptually simple
 - No dynamic allocation
- Most of the code is divided into tasks
- SPI bus used to communicate with Android
- UART bus used for logs and minimalistic console

EC Tasks

- idle ➔ system events, timers
- hook
- nugget ➔ system control task
- AVB ➔ secure boot management
- faceauth ➔ biometric data
- identity ➔ identity documents support
- keymaster ➔ key generation and cryptographic operations
- weaver ➔ storage of secret tokens
- console ➔ debug terminal and logs

Communication with the chip



Firmware security

- No dynamic allocation → no UaF and similar
- Secure boot (images are signed and verified at boot)
- MPU to give permissions to the memory partitions
 - Custom interface to set the eXecute permission
 - No WX permissions by default
- Only software protection: hardcoded stack canary

What we show today

- **Fuzzing** is useful also against Titan M
 - Even on such contrainted target, we can get interesting results
- Two approaches
 - Black-box fuzzer vs emulation-based fuzzer
- Exploiting **without** debuggers or stack traces
- How a single **software** vulnerability can lead to
 - Code execution
 - Compromise of the security properties guaranteed by the chip

Blackbox fuzzing

Black Box fuzzing

- Target: tasks
- Arbitrary messages with nosclient
 - Known format of the messages
 - We get a return code, and an actual response if successful
- **Mutate** the message, **check** return code
 - If greater than 1, *something* happened

external/nos/host/generic/nugget/include/application.h

```
enum app_status {
    /* A few values are common to all applications */
    APP_SUCCESS = 0,
    APP_ERROR_BOGUS_ARGS, /* caller being stupid */
    APP_ERROR_INTERNAL,   /* application being stupid */
    APP_ERROR_TOO_MUCH,   /* caller sent too much data */
    APP_ERROR_IO,          /* problem sending or receiving data */
    APP_ERROR_RPC,         /* problem during RPC communication */
    APP_ERROR_CHECKSUM,   /* checksum failed, only used within protocol */
    APP_ERROR_BUSY,        /* the app is already working on a command */
    APP_ERROR_TIMEOUT,     /* the app took too long to respond */
    APP_ERROR_NOT_READY,   /* some required condition is not satisfied */
}
```

Implementation

- Plug libprotobuf-mutator¹ in nosclient
 - Very straightforward
 - `void Mutate(protobuf::Message* message, size_t max_size_hint);`
- Basic corpus generation
 - Messages are quite simple
 - Start from empty ones, but add some non-trivial fields
- Store and triage inputs generating faulty states

Results

Firmware: 2020-09-25, 0.0.3/brick_v0.0.8232-b1e3ea340

- 2 buffer overflows (1 exploited for code exec)
- 4 null pointer dereferences
- 2 unknown bugs causing a reboot

Firmware: latest (at the time), 0.0.3/brick_v0.0.8292-b3875afe2

- 2 null-ptr deref still make the chip crash
- Bug reported → not a vulnerability

All of this after a few minutes of fuzzing...

Comments and limitations

- ✓ Bugs!
- ✓ Very simple to implement
- ✓ Decent performance: ~74 msg/sec
- ✓ Testing in real world

- ✗ Only “scratching the surface”
- ✗ Prone to false positives
- ✗ Detection is tricky
- ✗ Risk of bricks

Bottom line: hard to know what's going on the target

Emulation-based fuzzing

Switching to emulation-based

- We know how the OS works
- We can leak arbitrary memory with an exploit on an old firmware
 - Helps setting up memory
- With emulation, we control what is executed
 - Good feedback for a fuzzer

Emulating Titan M

- Played with several frameworks
- Choice: Unicorn²
- Why?
 - Emulates CPU only
 - We do not care about full-system emulation
 - Easy to setup & tweak
 - Integrates nicely with AFL++



Fuzzing with AFL++

- AFL++ in Unicorn mode
 - Instrument anything that can be emulated with Unicorn
 - Fuzz with the classic AFL experience
- Once the emulator works, not much needs to be done
 - place_input_callback to copy input sample
 - Crashes detected at Unicorn errors (e.g. UC_ERR_WRITE_UNMAPPED)
- Custom mutators depending on needs
 - AFL_CUSTOM_MUTATOR_LIBRARY=<mutator.so>
 - AFL_CUSTOM_MUTATOR_ONLY=1 to use only that one

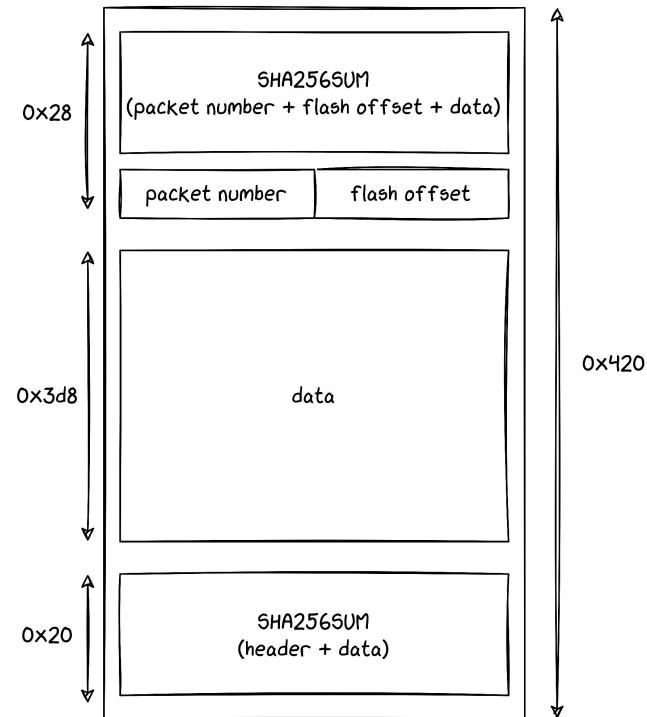


What to fuzz

- Pretty much anything!
- All you need is:
 - An entry point
 - Valid memory state
 - Registers set at the right values
 - One or more exit points
- Keep attack surface into account

SPI rescue feature

- SPI rescue allows to flash new firmware
 - No password required
 - Wipes user data
 - Can be triggered from bootloader
- Firmware sent as rec file



The Boot ROM

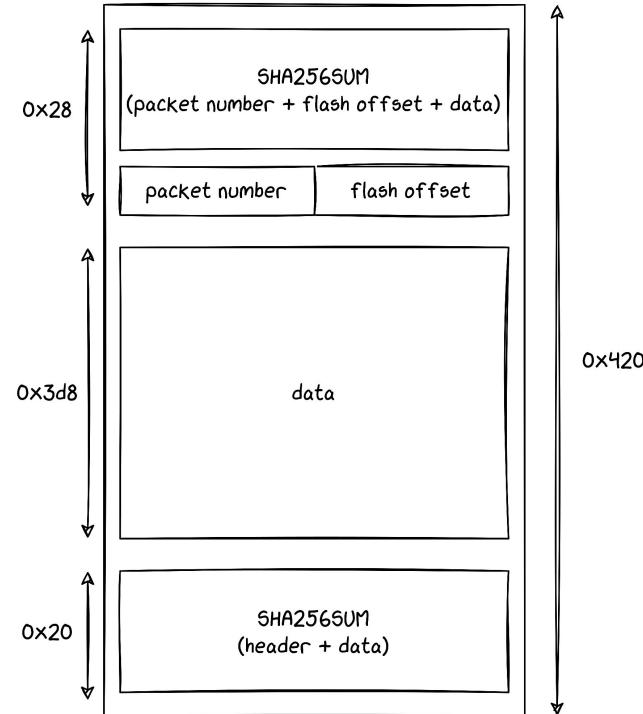
- Thanks to the 1-day exploit, we leaked the Boot ROM
- A bug there would be disastrous
- Not much code to test (only 16 KB)
- Idea: fuzz the image loader
 - We could flash them with SPI rescue

... no interesting results

- The function is simple, and not processing much
- Samples are just image headers

SPI rescue handler

- Focus on the rescue feature itself
- Are input files parsed and processed correctly?
- This time input is structured
 - Let's mutate it smartly :)
- We use FormatFuzzer⁴
 - Allows to generate and parse binary files
 - Follows the bt template format, from the O10 editor
 - Requires a modified version of AFL++



... also this time, no bugs (but some interesting internals revealed)

Going back to the tasks

- Tasks use protobuf
- Rely again libprotobuf-mutator
 - With some tricks to embed the message name in the bytes it generates
- Focused on Identity and Keymaster
 - The largest and most complex tasks
 - We fuzzed Weaver too, but it is not as interesting
- First, can we find the same bugs we know about?

Yes! (apart from one...)

There is no free lunch

- Emulation is not a silver bullet!
- Embedded targets → hw-dependant code everywhere...
 - Lots of hooks
 - Code that can't be exercised
 - Especially true in system functions
- A bug doesn't always make Unicorn crash
 - No ASAN-like instrumentation
 - In-page overflows, off-by-ones won't be detected
- No full system emulation → miss some parts of code
 - No system state
 - The bug we missed makes the scheduler crash
 - ... and we don't emulate the scheduler 😞

Tweaks

- Much more capabilities compared to pure black-box
- A few heuristics we implemented:
 - Monitor consecutive reads in the Boot ROM → spot buggy `memcpy`
 - Hook accesses to specific global buffers
 - Even more specific ones on different commands
- At the same time, everything comes at a cost
 - Hooks impact performance
 - In our case, not a big deal due to very specific harnesses

The vulnerability

CVE-2022-20233

- param_find_digests_internal
 - Checks digest tags in KeyParameter objects
- Out-of-bounds write of 1 byte to 0x1
 - Can be repeated multiple times
 - Huge constraints on the offset
- Looks like a minor issue...

```
message KeyParameter {  
    Tag tag = 1;  
    uint32 integer = 2;  
    uint64 long_integer = 3;  
    bytes blob = 4;  
}  
  
message KeyParameters {  
    repeated KeyParameter params = 1;  
}
```

CVE-2022-20233

```
ldr.w  r1,[r2,#-0x4]
ldr    r3,[PTR_DAT_0005d808] ; 0x20005
cmp    r1,r3
bne    increment_loop_vars
ldr    r3,[r2,#0x0]
uxtb   r0,r3
cmp    r0,#0x4
bhi    error_exit
movs   r1,#0x1
lsl.w  r0,r1,r0
tst    r0,#0x15
beq    error_exit
strb   r1,[r7,r3]
```

```
if (((nugget_app_keymaster_KeyParameter *)(offset + -1))->tag ==
     0x20005) {
    masked = *offset & 0xff;
    if ((4 < masked) || ((1 << masked & 0x15U) == 0)) {
        return 0x26;
    }
    *(undefined *) (buffer + *offset) = 1;
    *param_3 = *param_3 + 1;
    *param_4 = offset;
}
```

CVE-2022-20233

```
ldr.w  r1,[r2,#-0x4]
ldr    r3,[PTR_DAT_0005d808] ; 0x20005
cmp    r1,r3
bne    increment_loop_vars
ldr    r3,[r2,#0x0]
uxtb   r0,r3
cmp    r0,#0x4
bhi    error_exit
movs   r1,#0x1
lsl.w  r0,r1,r0
tst    r0,#0x15
beq    error_exit
strb   r1,[r7,r3]
```

```
if (((nugget_app_keymaster_KeyParameter *)(offset + -1))->tag ==
     0x20005) {
    masked = *offset & 0xff;
    if ((4 < masked) || ((1 << masked & 0x15U) == 0)) {
        return 0x26;
    }
    *(undefined *) (buffer + *offset) = 1;
    *param_3 = *param_3 + 1;
    *param_4 = offset;
}
```

CVE-2022-20233

```
ldr.w  r1,[r2,#-0x4]
ldr    r3,[PTR_DAT_0005d808] ; 0x20005
cmp    r1,r3
bne    increment_loop_vars
ldr    r3,[r2,#0x0]
uxtb   r0,r3
cmp    r0,#0x4
bhi    error_exit
movs   r1,#0x1
lsl.w  r0,r1,r0
tst    r0,#0x15
beq    error_exit
strb   r1,[r7,r3]
```

```
if (((nugget_app_keymaster_KeyParameter *)(offset + -1))->tag ==
     0x20005) {
    masked = *offset & 0xff;
    if ((4 < masked) || ((1 << masked & 0x15U) == 0)) {
        return 0x26;
    }
    *(undefined *)(buffer + *offset) = 1;
    *param_3 = *param_3 + 1;
    *param_4 = offset;
}
```

0xdeadbeef

CVE-2022-20233

```
ldr.w  r1,[r2,#-0x4]
ldr    r3,[PTR_DAT_0005d808] ; 0x20005
cmp    r1,r3
bne    increment_loop_vars
ldr    r3,[r2,#0x0]
uxtb   r0,r3
cmp    r0,#0x4
bhi    error_exit
movs   r1,#0x1
lsl.w  r0,r1,r0
tst    r0,#0x15
beq    error_exit
strb   r1,[r7,r3]
```

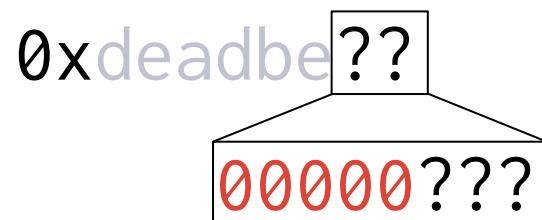
```
if (((nugget_app_keymaster_KeyParameter *)(offset + -1))->tag ==
0x20005) {
masked = *offset & 0xff;
if ((4 < masked) || ((1 << masked & 0x15U) == 0)) {
    return 0x26;
}
*(undefined *)(buffer + *offset) = 1;
*param_3 = *param_3 + 1;
*param_4 = offset;
}
```

0xdeadbeef

CVE-2022-20233

```
ldr.w  r1,[r2,#-0x4]
ldr    r3,[PTR_DAT_0005d808] ; 0x20005
cmp    r1,r3
bne    increment_loop_vars
ldr    r3,[r2,#0x0]
uxtb   r0,r3
cmp   r0,#0x4
bhi   error_exit
movs   r1,#0x1
lsl.w  r0,r1,r0
tst   r0,#0x15
beq   error_exit
strb  r1,[r7,r3]
```

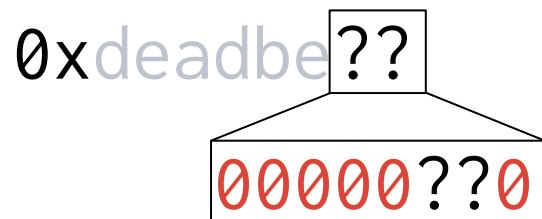
```
if (((nugget_app_keymaster_KeyParameter *)(offset + -1))->tag ==
     0x20005) {
    masked = *offset & 0xff;
    if ((4 < masked) || ((1 << masked & 0x15U) == 0)) {
        return 0x26;
    }
    *(undefined *) (buffer + *offset) = 1;
    *param_3 = *param_3 + 1;
    *param_4 = offset;
}
```



CVE-2022-20233

```
ldr.w  r1,[r2,#-0x4]
ldr    r3,[PTR_DAT_0005d808] ; 0x20005
cmp    r1,r3
bne    increment_loop_vars
ldr    r3,[r2,#0x0]
uxtb   r0,r3
cmp    r0,#0x4
bhi    error_exit
movs   r1,#0x1
lsl.w  r0,r1,r0
tst    r0,#0x15
beq    error_exit
strb   r1,[r7,r3]
```

```
if (((nugget_app_keymaster_KeyParameter *)(offset + -1))->tag ==
     0x20005) {
    masked = *offset & 0xff;
    if ((4 < masked) || ((1 << masked & 0x15U) == 0)) {
        return 0x26;
    }
    *(undefined *) (buffer + *offset) = 1;
    *param_3 = *param_3 + 1;
    *param_4 = offset;
}
```



What can we do?

- Multiple ways to reach the vulnerable code
 - A few different command handlers call it
 - Different base addresses for the OOB-write
- Titan M's memory is completely static
 - All structures are always located at the same addresses
- Setting one byte can be enough to break the system

Our approach

- Generate all writable addresses
- Highlight them in Ghidra
- ...

KEYMASTER_SPI_DATA

c8	92	01	00

`void * callback_addr`

`char * cmd_request_addr`

`char * cmd_response_addr`

Our approach

- Generate all writable addresses
- Highlight them in Ghidra
- ...

KEYMASTER_SPI_DATA

c8	01	01	00

`void * callback_addr`

`char * cmd_request_addr`

`char * cmd_response_addr`

What to overwrite

KEYMASTER_SPI_DATA

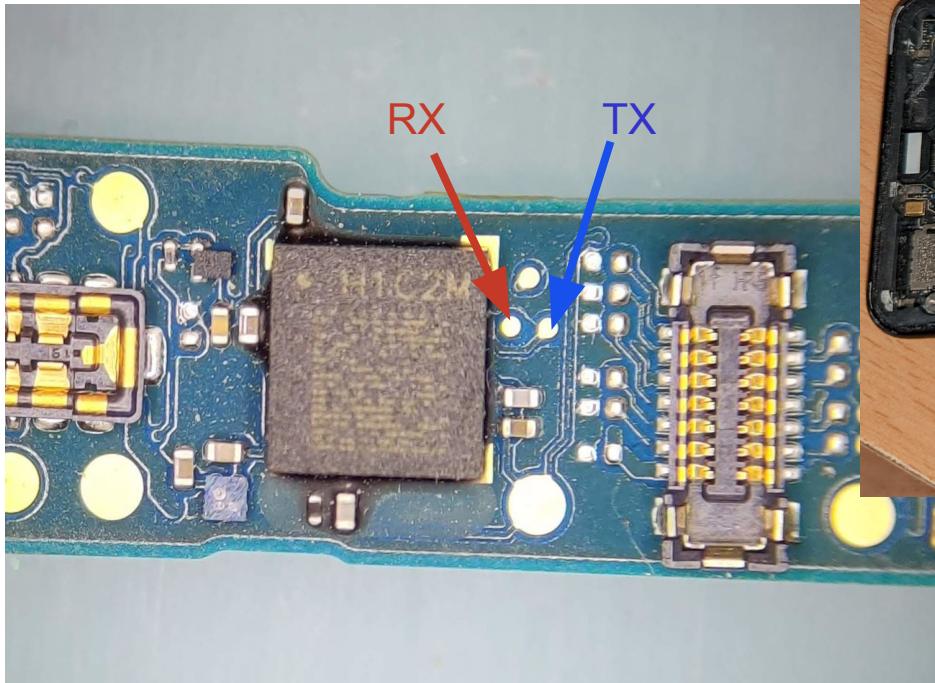
- Global structure
- Stores info about SPI commands
- cmd_request_addr: where to store incoming Keymaster requests
- 0x192c8 → 0x101c8



But first...

- Remainder:
 - Communication through nosclient
 - Send request using Android libs
 - Get a return code and (maybe) a response
 - A few logs on logcat
- What if we crash the chip?
 - Error code 2
- That's it
- Debugging an exploit is... challenging

Accessing the UART



UART console

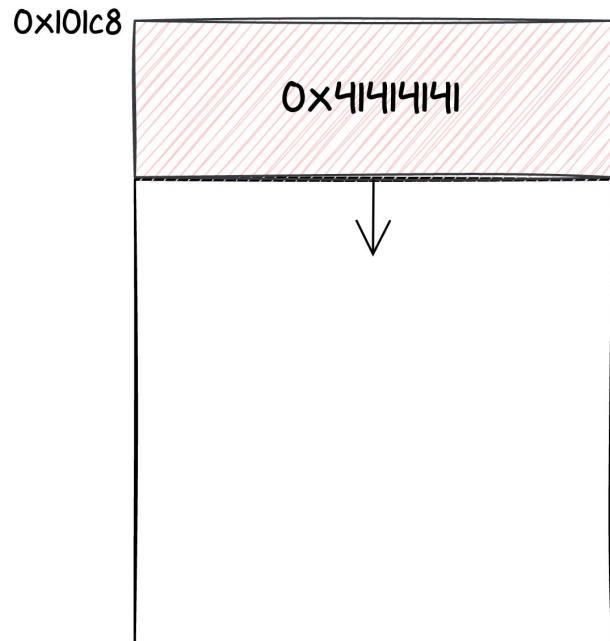


```
$ picocom /dev/ttyUSB0 -b 115200
[Image: RW_A, 0.0.3/chunk_ab7976980-a9084b7 2021-12-07
18:40:23 android-build]
[1.694592 Inits done]
[1.695460 update_rollback_mask: stop at 1]
[1.695884 gpio_wiggling: AP_EL2_LOW_IRQ = 0]
Console is enabled; type HELP for help.
>
> help
Known commands:
    apfastboot      history      repo        taskinfo      version
    board_id        idle         sleepmask   timerinfo
    help            reboot       stats        trngstats
HELP LIST = more info; HELP CMD = help on CMD.
```

- Allows basic interaction
- Prints logs
 - Useful when exploiting

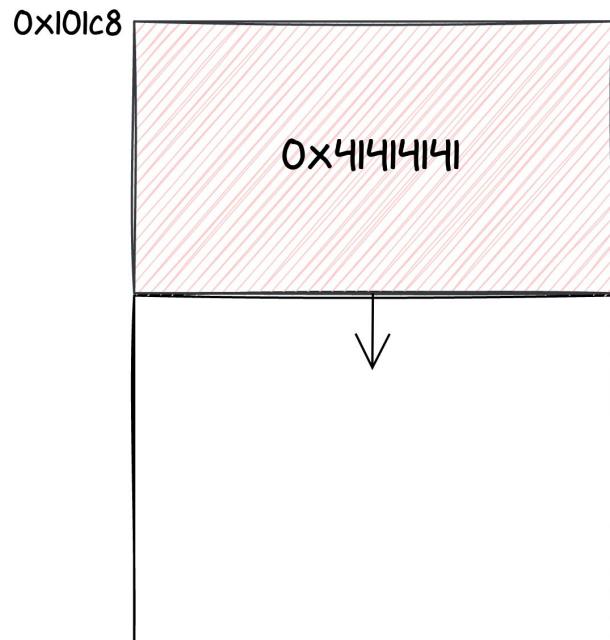
So, what's in 0x101c8?

- Data doesn't seem to be used
- How do we hijack execution flow?
- Idea:
 - Send progressively bigger payloads
 - In parallel monitor the UART
 - ... and see what happens



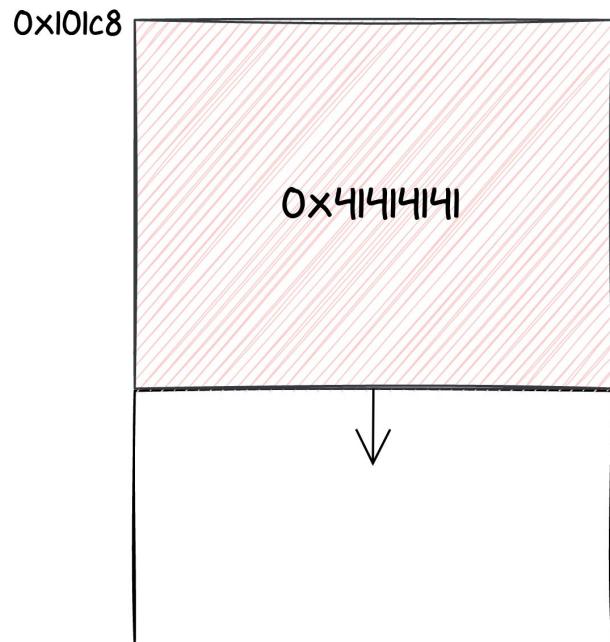
So, what's in 0x101c8?

- Data doesn't seem to be used
- How do we hijack execution flow?
- Idea:
 - Send progressively bigger payloads
 - In parallel monitor the UART
 - ... and see what happens



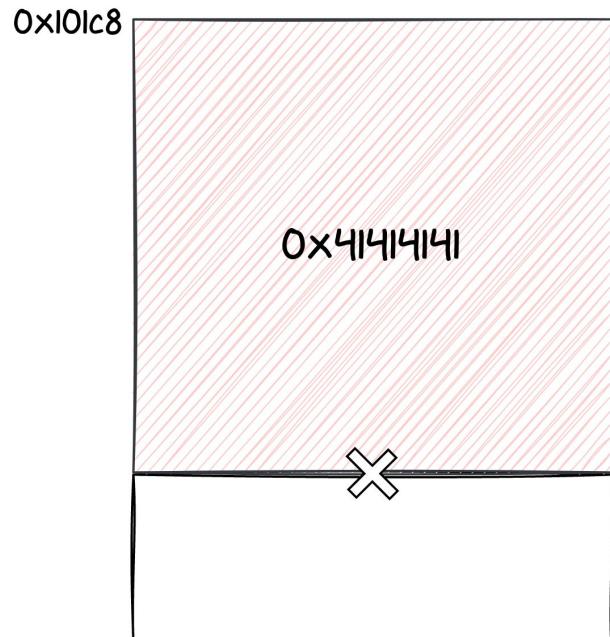
So, what's in 0x101c8?

- Data doesn't seem to be used
- How do we hijack execution flow?
- Idea:
 - Send progressively bigger payloads
 - In parallel monitor the UART
 - ... and see what happens



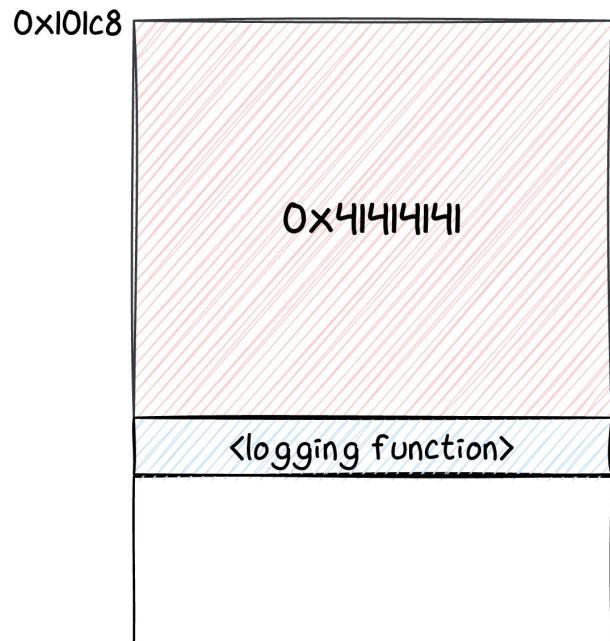
So, what's in 0x101c8?

- Data doesn't seem to be used
- How do we hijack execution flow?
- Idea:
 - Send progressively bigger payloads
 - In parallel monitor the UART
 - ... and see what happens
- At some point, the chip starts crashing



So, what's in 0x101c8?

- Data doesn't seem to be used
- How do we hijack execution flow?
- Idea:
 - Send progressively bigger payloads
 - In parallel monitor the UART
 - ... and see what happens
- At some point, the chip starts crashing
- What if we put a valid address at the end?



So, what's in 0x101c8?

- Data does
- How does
- Idea:
 - Sends
 - In pa
 - ... an
- At some
- What if w
- \o/

The screenshot shows a terminal window with a black background and white text. At the top, it says "0x101c8". Below that, there is a large amount of log output starting with "--- UART initialized after reboot ---" and ending with a task list. The task list has columns for Task ID, Ready, Name, Events, Time (s), StkUsed, and Flags. The tasks listed are:

Task	Ready	Name	Events	Time (s)	StkUsed	Flags
0	R << idle >>		00000000	0.000116	80/ 512	0000
1	R	HOOKS	20000000	0.001708	152/ 640	0000
2		NUGGET	00000000	0.012108	168/1024	0000
3		FACEAUTH	00000000	0.000456	80/2048	0000
4		AVB	00000000	0.004440	88/4096	0000
5		KEYMASTER	00000000	0.015640	88/9600	0000
6		IDENTITY	00000000	0.000220	88/1952	0000
7		WEAVER	00000000	0.010372	240/1024	0000
8		CONSOLE	00000000	0.024296	80/ 576	0000

On the right side of the terminal window, there is a large redacted area with the text "0x41414141" written vertically. Below this, a blue horizontal bar contains the text "gging function>".

Exploiting

- Our guess:
 - We are actually in the stack of a task (idle)
 - We overwrite a function pointer that was pushed to the stack
 - At some point, the function jumps back to it
- From here on, things get complex
 - No space to write a ROP chain there
 - We need to move \$sp
- In the end, we send another command to complete the exploit
- Blogpost arriving soon :)

Impact

- Control the execution flow of the chip
 - We are not able to reconfigure the MPU
 - ... but we can do pretty much anything using ROP
- We implemented again a leak command
 - This time based on a 0-day
 - Data is not erased by the downgrade like before!
 - We can leak all the secrets stored in the chip's memory

```
sargo:/data/local/tmp # ./nosclient leak 0x0 0x10
00 00 02 00 99 14 00 00 b9 3e 00 00 b9 3e 00 00
```

Can we leak strongbox keys?

StrongBox

- StrongBox: hardware-backed version of Keystore
 - Generate, use and encrypt cryptographic material
- Titan M does not store keys
 - Key blobs encrypted with a **Key Encryption Key**
 - This KEK is derived in the chip from various internal elements
 - Key blobs are sent to the chip to perform crypto operations
 - root can **use** any key, but not **extract** it

Strongbox

KEKs are derived from a **key ladder**

- Still quite mysterious since we did not reverse it
 - It uses
 - An internal *root key*
 - Not readable from the Titan M firmware
 - A *Root Of Trust* provided by the bootloader at first boot
 - A *salt* that is randomly generated when RoT is provisioned
- We can leak most of the secrets, but not the key ladder root key

Strongbox

There are 3 commands to use strongbox keys:

- *BeginOperation*
 - Contains the keyblob and the characteristics of the key
 - The chip will decrypt the keyblob
 - And save it for later into a **fixed address**
- *UpdateOperation*
 - Contains the data on which the operation is performed
 - Return the output bytes
- *FinishOperation*
 - Contains the data on which the operation is performed
 - Return the output bytes
 - End the operation

Leak strongbox keys

Our strategy:

1. Get the keyblob from the device
 - Stored in /data/misc/keystore/persistent.sqlite
2. Forge a *BeginOperation* request
3. Leak the decrypted key from the chip memory

```
sargo:/data/local/tmp # ./nosclient leak_kb  
- Key name: strongbox (size: 128)  
sargo:/data/local/tmp # ./nosclient leak_kb -k strongbox  
f3 7d 3d 7d ca 56 5e a0 18 ff 83 76 19 39 eb c1
```

Conditions

- Ability to send commands to the chip
 - Being root
 - Or direct access to the SPI bus
- Access to the keyblobs
 - Being root
 - Or find a way to bypass FBE...

Mitigation

```
KeyGenParameterSpec spec = new KeyGenParameterSpec.Builder("key_name",  
    KeyProperties.PURPOSE_ENCRYPT | KeyProperties.PURPOSE_DECRYPT)  
    .setBlockModes(KeyProperties.BLOCK_MODE_CBC)  
    .setEncryptionPaddings(KeyProperties.ENCRYPTION_PADDING_PKCS7)  
    .setIsStrongBoxBacked(true)  
    .setUserAuthenticationRequired(true)
```

Conclusion

Conclusion

- Fuzzing allowed us to find more bugs
- With black box, you easily get the surface bugs
- Emulation-based fuzzing is particularly effective of such target
 - Yet few tricks are required to optimize the results
- We found a critical 0-day
 - Allowed us to execute code on the chip
 - Permit to leak anything from the chip's memory
- A single software vulnerability is enough to leak strongbox keys

Tools & resources:

<https://github.com/quarkslab/titanm>

Thank you!

contact@quarkslab.com

Quarkslab



@max_r_b
@DamianoMelotti